

Testing data augmentation techniques and VGG convolutional neural networks for multiclass image classification on the CIFAR-10 dataset

Julian Cabezas Pena Student ID: a1785086
Assignment 2. Deep Learning Fundamentals, 2020
University of Adelaide
`julian.cabezaspena@student.adelaide.edu.au`

Abstract

In the last years, the field of image classification has been heavily influenced by the use of convolutional neural networks (CNN), that use convolution to extract data from a small portion of the image, usually called kernel, reducing the number of parameters compared with traditional neural networks and archiving better performance. One of the key aspects of the application of a CNN is the inclusion of data augmentation techniques, that can help reduce overfitting and increase performance, especially in small datasets. The VGG architecture of CNN appeared in 2015 and reached the state of the art in the ILSVRC challenge, using 3x3 kernels and from 11 to 19 convolutional layers. In this research, different data augmentation methods were tested using two different versions of the VGG architecture, with the objective of determining whether the inclusion of data augmentation techniques can have a larger or smaller effect that the increase on the deepness of the neural network. For this purpose, the CIFAR-10 dataset was used to test different data augmentation and architectures. The results indicated that the data augmentation methods can have a significant impact in the performance of this model, archiving the best performance using a random horizontal flip together with a random crop of the image, using the VGG19 architecture. The final model reached a accuracy of 88.82% on the test set, outperforming other implementation that do not include data augmentation but being far from the state of the art, dominated by transfer learning models. This research showed that practitioners should invest significant attention into the application of data augmentation techniques.

1. Introduction

In the field of machine learning and pattern recognition, classification problems are one of the most common tasks. In these kind of problems, an algorithm takes a vector of

data, and assigns it to one or many discrete classes [2]. One of the fields where classification algorithms are extensively used is the field of image classification, where rapid advances in the field have been driven by an accelerated development of the web, that each day generates an immense amount of image data. The field of image classification aims to organize image data into a limited amount of classes, based on the characteristics or features of the image. [18]

Early approaches to image classification include the Perceptron, that corresponded to a binary linear classifier [12], or, in the beginning of the century, the use of Support Vector Machine (SVN) [17] to address non-linear boundaries between classes. Another approach to this problem is the use of Artificial Neural Networks (ANN) of various kinds [10]. Although these algorithms can produce relatively good results depending on the dataset, in the beginning of the 2010 decade, the paradigm of image classification shifted towards the use of deep convolution networks, that could in some occasions halve the error rate of previous methods [10].

A convolution neural network is typically defined as a neural network that presents one or more convolutional layers. A convolutional layer is a neuron that covers an adjacent portion of an image, typically groups of 3x3 or 5x5 pixels, transforming the image using a shared set of weights. This technique can cause a major reduction of the number of parameters when comparing it with fully connected neural networks [1]. In general terms, the convolutional neural networks consists in a series of convolution layers followed by activation functions, while between the convolutional layers pooling layers are used. These convolution blocks are then followed by series of fully connected layers, ending with a soft-max function used to predict the target class [10].

The inclusion of the convolutional layers made a variety of new architectures of neural networks appear. One of the most popular architectures of deep convolutional networks is the VGG family of networks (named after the Visual Geometry Group of the University of Oxford), that introduced

the use of a series of very small (3x3 pixels) convolutional layers, archiving the state of the art in the ImageNet Large-Scale Visual Recognition Challenge (ILSVRC) of 2015 [15]

Although the complexity and deepness of the neural network is usually a big factor in the performance of the neural network [5], the training of these deep networks usually can produce overfitting, and relies on big quantities of data such as ImageNet, while in many domains of applications this data is not available. To overcome these problems, different data augmentation techniques can be used. These techniques aim to enhance the quality or size of the data contained in the image, partially solving the limited data problem [14]

The objective of this research is to implement two different VGG architectures with different number of layers (VGG11, 11 layers and VGG19, 19 layers) to perform multi class image classification using the CIFAR-10 dataset as study case. These architectures will be tested using combinations of easily available and common data augmentation methods to study how these techniques can increase performance and decrease overfitting. A secondary objective of this study is to determine whether the number of layers of a neural network can be more or less important than the correct application of data augmentation techniques.

2. Background

As the CIFAR-10 is a relatively known database, multiple neural network architectures has been tested on it. As an example, Liu and Deng [11] applied the VGG16 neural network architecture to the above mentioned dataset, adapting the fully connected layers to match the size (32x32) of the images of this dataset, while at the same time increasing the dropout rate to avoid overfitting, archiving 78.29%, the same network. When the authors included batch normalization, they archived an accuracy of 91.55%.

The current (2020) state-of-the-art for image classification on this dataset is held by Google Research, that implemented a transfer model called Big Transfer (BiT), accomplishing a 99.4% of accuracy. This algorithm consists in the application of a pre-trained model (BiT-L), that was trained on the JFT-300M dataset, consisting in 300 million labelled images, and then applied on the CIFAR-10 dataset for fine-tuning [8]

The most closely related competitors of the VGG could be AlexNet and ResNet, that were developed in 2012 and 2015 respectively. Alexnet is a neural network that consists in eight layers, with the first five being convolutional layers with kernel sizes of 11x11 (first layer), 5x5 (second layer), and 3x3 (third to fifth layer), that are followed by max pulling layers in the case of the first, second and fifth convolutions. Each convolution is activated using a ReLU function, that proved more effective than a sigmoid or tanh function. This neural network finishes with three fully con-

nected layers, that output the predicted class [9]. This neural network can be considered the precursor of the VGG architecture, that modified it by using a larger number of convolution layers of kernel size 3x3, and the dropout regularization to avoid overfitting.

On the other hand, the residual networks (ResNet) are the architectures that came after VGG, managing to solve the gradient vanishing problem. The gradient vanishing is a recurrent issue in deep neural networks, that refers to the problem that occurs when, thought the process of gradient descent, the weight updates became smaller as the layer is farther from the output layer, due to the repeated derivation of the functions that are involved in the process of backpropagation [1, 16]. To approach this phenomena, the ResNet presents the concept of residual learning, that implies a shortcut or skipping between blocks of layers (residual blocks), using a identity function, adding the output of the previous layer to the layer ahead of the block. In practice this means these neural networks can be trained using more layers than what previously presented, implementing ResNets of more than one thousand layers, pushing the limits of was considered depth in the CNN area, and avoiding the vanishing gradients problem [5, 6]

Even though the VGG results are not comparable to more modern neural networks, such as ResNet [5], or the above mentioned state of the art transfer learning techniques [8]. VGG network stand for its simplicity and for counting with variations of the same structure of different number of convolution layers, that can help approach the trade-off between data augmentation and number of layers, that is one of the objectives of this study.

3. Methods

3.1. CIFAR-10 Dataset

The CIFAR-10 dataset (named after the Canadian Institute of Advanced Research), consists in 60000 images of small size, retrieved from online search (using Google or Altavista) and resampled to 32x32 pixels resolution. These images were manually picked to only consider photo-realistic images, containing only one prominent instance of the object being classified. These images were manually labelled into 10 mutually exclusive classes [10].

An image example for each of the ten classes that are represented in this dataset is shown in Figure 1:

The database contains 6000 images of each class. Of the total database, 50000 images (5000 for each class) were considered for the train dataset and the remaining 10000 for the test dataset (1000 for each class)

3.2. VGG11 and VGG19 neural networks

The VGG architectures that were used in this study were introduced in 2015 by the Visual Geometry Group of the

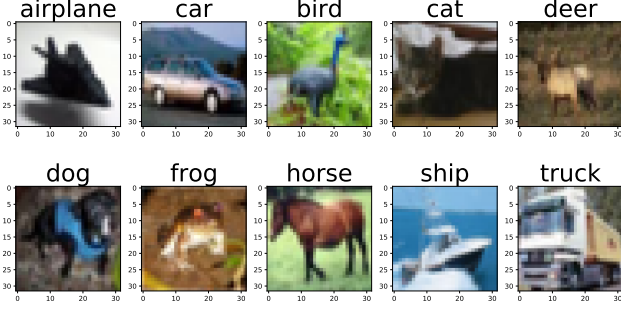


Figure 1. Samples of the ten classes in CIFAR-10 dataset

University of Oxford, and was designed to classify the ImageNet dataset. In these architectures, the images are passed to several stacks of convolutional layers of 3x3 kernel size, with a stride of 1 pixel, causing that the spatial resolution of the image is preserved, each convolution layer is followed by a ReLU activation function. After each stack of convolutional layers, a maximum pooling layer with a window size of 2x2 and a stride of 2 is used [15], that implies that in practice after each max pooling layer the spatial resolution of the image is halved.

Finally, after all the convolution layers are applied, three linear fully connected layers are applied, using a dropout regularization process in the first two fully connected layers (in the training process). The last of the fully connected layers outputs a number of values equal to the number of classes (in this case 10), finalizing with a Softmax function [15], that is used to determine the predicted value selecting the maximum output value.

The two neural networks that were used are the VGG11 and the VGG19, named after the number of layers in their structure. In Figure 2, it is possible to appreciate that the two VGG architectures share a common structure, with the VGG19 having 8 extra convolution layers. To adapt the original VGG architecture to the CIFAR-10 dataset, the input and output sizes of the fully connected layers were modified.

3.3. Loss function and optimization criterion

To get the predicted value of the neural network, a Softmax function, that takes the last output of the last neural network, was used:

$$\text{Softmax}(x_i) = \frac{\exp(f(x_i))}{\sum_{j=1}^n \exp(f(x_j))} \quad (1)$$

where x_i is the i output element result of the last fully connected layer of the VGG neural network. As the function includes a division by the sum of the exponentials of the outputs, the sum the Softmax output values is always 1 [16].

To calculate the loss in each batch, the cross-entropy function for a set of parameters θ is used [3]:

$$\text{Loss}(\theta) = - \sum_{i=1}^n y_i \log(f(x_i)) \quad (2)$$

where y_i is the one-hot encoded vector of true labels.

This loss function was optimized using a mini batch Stochastic Gradient Descent, as in the original VGG publication [15], while in this case a batch size of 10 was used. The weights of the convolutional layers were initialized using the Kaiming method, that has proven to be useful when dealing with non linear activation functions such as ReLU [4]

3.4. Data Augmentation techniques

As data augmentation techniques are described in literature as very important to avoid overfitting and to take advantage of small datasets [14], such as CIFAR-10, three data augmentation techniques and its combinations were tested on the VGG11 and VGG19 architectures:

Random Horizontal Flip: This simple flipping data augmentation technique consists in a 180 degree flipping over the horizontal axis of the image., It has proven effective in ImageNet and CIFAR-10 dataset [14], and was used by Simonyan and Zisserman [15] in the original formulation of the VGG architecture.

Random Clip: Random cropping or clipping can be a very effective way to simulate the effect of a translation or a feature not being in the centre [14], in this case a random padding with a maximum of 4 pixels was used, and then the image was clipped to preserve the dimensions of the image (32x32), filling the padded pixels with zero values.

Color Jitter: A colour space transformation was applied to simulate different conditions of brightness and contrast of the image [14], this tool was configured to modify the brightness, contrast and saturation of the image, with multipliers ranging from 0.5 and 1.5.

Examples of the application of these data augmentation methods can be found in Figure 3, where these methods are shown using samples of the CIFAR-10 dataset.

The only data augmentation that was applied to all the simulations is the normalization of the values of the images using $mean = 0.5$ and $sd = 0.5$, that was applied after the transformation, making the values range from -1 to +1 [16].

3.5. Hyperparameter tuning

In order to select the best combinations of data augmentation methods and number of convolution layers (VGG11 and VGG19), all the possible combinations of the above mentioned data augmentation methods were used to train the VGG architectures in the train set, calculating the accuracy on the train and validation set for each of the epochs.




Figure 2. VGG11 (top) and VGG19 (bottom) architectures used in this study, the convolution layers are denoted as conv(kernel size) - (output channels), each convolution layer is followed by a ReLu activation function [15]




Figure 3. Samples of the data augmentation methods utilized in this study

As the number of combinations is high (8 for each architecture), only 50 epochs were used, and the best combination was selected, using the from the maximum accuracy in the validation set.

$$Accuracy = \frac{Number\ of\ correct\ predictions}{Total\ number\ of\ predictions} \quad (3)$$

The selected combination of methods were then trained using 100 epochs to test three different learning rates (0.01, 0.05 and 0.001), also selecting this parameter using the higher accuracy in the validation set

The final combination of hyperparameters was tested on the totality of the training data using 100 epochs, and the resulting model was used to predict on the test data, calculating the overall accuracy of the model and the accuracy of each class

4. Code and processing

The code, along with the requirements and setup instructions for the reproduction of the above-mentioned method, can be found in the following GitHub repository: https://github.com/juliancabezas/Convolutional_Neural_Networks_CIFAR10

All the processing was performed in a Linux OS with NVIDIA GeForce GTX 1050 with CUDA support, and at the same time using the Google Colab platform (<https://colab.research.google.com/>) free tier GPU processing capabilities, that include a NVIDIA K80 graphical unit.




Figure 4. Training accuracy (left), validation accuracy (centre) and cross-entropy loss (right) of the VGG11 neural network using different data augmentation techniques. NO AUG: No data augmentation, CJ: Color Jitter, RH: Random Horizontal Flip, RC: Random Crop

5. Results and discussion

5.1. Data augmentation testing

In Figure 4 it is possible to appreciate that the VGG11 implementation without data augmentation tends to archive maximum training accuracy early, archiving 99.07% accuracy in epoch number 36, while at the same time presenting a validation accuracy of 76.87%, that is close to the maximum validation accuracy (78.26% in epoch 49). On the other side, all the combinations of data augmentation techniques excepting the color jitter transformation helped decrease this effect, while at the same time greatly improving the validation accuracy. The best results are produced with the combination of all three data augmentation techniques, presenting a maximum validation accuracy of 84.99%, and with the combination of random horizontal flip (RH) and random crop (RC), with a maximum validation accuracy of 85.56%. Showing that in this case, the data augmentation can cause a dramatic increase in performance.

The behaviour of the accuracy and loss in the case of the VGG19 architecture (Figure 5) follow a similar pattern as in the VGG11 implementation, with the version of the neural network without data augmentation reaching 99.08% training accuracy in epoch 40, and a maximum validation accuracy of 81.12%, and showing better validation performance than the one with the color jitter transformation. In this case, the best performance is also accomplished by the combination of the random horizontal flip and the random crop, archiving 86.36% validation accuracy in epoch 17

When comparing the VGG11 and VGG19 implementations without extra data augmentation techniques, it is possible to notice that the increase of convolutional layers generated an increase of around 3% in the maximum valida-

tion accuracy (78.26% in VGG11 and 81.12% in VGG19). On the other side, when comparing both neural networks with the random crop and random horizontal flip augmentations, the performances of both algorithms are close, with less than 1% difference in accuracy (85.56% in VGG11 and 86.36% in VGG19). These results show that, in order to increase the performance of the models, choosing the correct data augmentation methods can produce more benefits than the increase of convolution layers, while at the same time being simpler, as the architecture of the network does not have to be modified.

As above mentioned, of the 16 combinations of architecture and data augmentation methods the best result was archived by the VGG19 neural network with random horizontal flip (RH) and random crop (RC) methods, that was chosen to continue with the next steps of the hyperparameter adjusting.

5.2. Learning rate

The results of the learning rate adjusting showed that in this particular case, the different learning rates did not produce great differences in performance or training time (Figure 6), with the best performance being archived in the epoch 98 (88.35% validation accuracy) using learning rate = 0.01.

5.3. Testing of the final model

The VGG19 architecture, trained using 100 epochs and a learning rate of 0.01 with random horizontal flip and random crop data augmentations methods, gave the test accuracy results showed in Table 1. It is possible to appreciate that some imbalance exists in the accuracy of the different classes, with the worst accuracies are presented in the Bird




Figure 5. Training accuracy (left), Validation accuracy (centre) and Cross-entropy loss (right) of the VGG19 neural network using different data augmentation techniques. NO AUG: No data augmentation, CJ: Color Jitter, RH: Random Horizontal Flip, RC: Random Crop

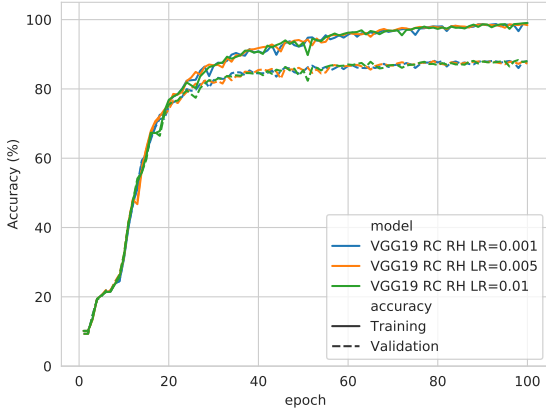


Figure 6. Learning rate tuning on the VGG19 with Random Horizontal Flip and Random Crop model. LR: learning rate

(78.60%) and Dog (77.40%). This results can be caused by the fact that bird image can be found both in the ground and the sky, making the variety of features in the bird images larger. The best accuracy in the Ship class (95.50%). The overall accuracy of the prediction on the test data was 88.82%.

5.4. Comparison with other implementations

Table 2 shows the comparison of the performance of the implemented model with other methods that has been tested on the same data. The top performance is observed in transfer methods (Bit-L and GPipe) [8, 7], that take advantage of other bigger dataset to pretrain the neural networks, only fine tuning the model parameters using the CIFAR-10 train-

Class	Test Accuracy
Airplane	91.00%
Car	94.30%
Bird	78.60%
Cat	82.10%
Deer	92.80%
Dog	77.40%
Frog	92.10%
Horse	91.20%
Ship	95.50%
Truck	93.20%
Overall accuracy	88.82%

Table 1. Model performance in the test data

ing data, reaching accuracies greater than 99%.

The trained VGG19 model performance is also outperformed by deeper ResNet models (ResNet-1001 and ResNet-164), that were also trained picking adequate data augmentation methods [6]. Although, on the other side, the presented model outperforms the results obtained by Sharma *et al* [13] using a 50 layer ResNet without data augmentation techniques, that accomplished only 78.10% accuracy, showing that even though this neural network counts with more that two times the layers of the VGG19 architectures, when it is applied without considering appropriate data augmentation techniques, can be outperformed by a simpler VGG neural network with good data augmentation. The implemented method also predictively outperforms older methods such as GoogleLeNet and Alexnet [13].

Method	Acc.	Reference
BiT-L	99.4%	Kolesnikov <i>et al</i> [8]
GPipe	99.4%	Huang <i>et al</i> [7]
ResNet-1001	95.38 %	He <i>et al</i> [6]
ResNet-164	94.54%	He <i>et al</i> [6]
VGG16 with BN	91.55%	Liu & Deng [11]
VGG19-RH-RC	88.82%	This Study
VGG16	78.29%	Liu & Deng [11]
ResNet50	78.10%	Sharma <i>et al</i> [13]
GoogleLeNet	71.67%	Sharma <i>et al</i> [13]
AlexNet	36.08%	Sharma <i>et al</i> [13]

Table 2. Performance comparison with other studies

6. Conclusion

This research showed the importance of data augmentation, that in this particular case caused an increase on performance that was larger than the inclusion of more convolutional layers. This insight can be particularly important for deep learning practitioners, showing the community that time invested in testing different data augmentation methods can produce great benefits and can sometimes be more effective than designing deeper neural networks, specially considering the fact that the implementation of data augmentation methods is usually straightforward in modern deep learning libraries.

The VGG architecture of neural networks, despite its relative oldness, was capable of producing relatively acceptable results when the adjusting of hyperparameters and preprocessing methods is well performed, showing the potential of simple architectures. Although transfer methods seem to greatly outperform single CNN implementations, showing the trend that deep learning could follow in terms of image classification in the future.

In future research, batch normalization as well as other commonly used data augmentation methods could be used to get better results, additionally, the normalization could be performing using the actual mean and standard deviation of the images in the dataset, possibly improving the performance of the VGG models. Additionally, the hypothesis that were investigated in this study can also be confirmed using other architectures of neural network that present different amount of layers, such as ResNet, and with a bigger variety of neural network architectures.

7. Bonus

This research provided insight over the importance of data augmentation methods, in this case comparing it with the increase of the number of layers in neural networks, using experimental results to back up this hypothesis.

References

- [1] Hamed Habibi Aghdam and Elnaz Jahani Heravi. *Guide to Convolutional Neural Networks: A Practical Application to Traffic-Sign Detection and Classification*. 2017.
- [2] Christopher M Bishop. *Pattern Recognition and Machine Learning*. Springer, New York, NY, USA, 2006.
- [3] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The Elements of Statistical Learning*. 2009.
- [4] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. In *Proceedings of the IEEE International Conference on Computer Vision*, volume 2015 Inter, pages 1026–1034, 2015.
- [5] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, 2016-Decem:770–778, 2016.
- [6] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Identity mappings in deep residual networks. *Computer Vision - ECCV 2016. ECCV 2016. Lecture Notes in Computer Science*, 9908:630–645, 2016.
- [7] Yanping Huang, Youlong Cheng, Ankur Bapna, Orhan Firat, Mia Xu Chen, Dehao Chen, HyounJoong Lee, Jiquan Ngiam, Quoc V. Le, Yonghui Wu, and Zhifeng Chen. GPipe: Efficient Training of Giant Neural Networks using Pipeline Parallelism. *arXiv e-prints*, page arXiv:1811.06965, 2019.
- [8] Alexander Kolesnikov, Lucas Beyer, Xiaohua Zhai, Joan Puigcerver, Jessica Yung, Sylvain Gelly, and Neil Houlsby. Big Transfer (BiT): General Visual Representation Learning. 2019.
- [9] Alex Krizhevsky. Learning Multiple Layers of Features from Tiny Images. Technical report, University of Toronto, Toronto, Canada, 2012.
- [10] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. ImageNet classification with deep convolutional neural networks. *Communications of the ACM*, 60(6):84–90, 2017.
- [11] Shuying Liu and Weihong Deng. Very deep convolutional neural network based image classification using small training sample size. *Proceedings - 3rd IAPR Asian Conference on Pattern Recognition, ACPR 2015*, pages 730–734, 2016.
- [12] F. Rosenblatt. The Perceptron - A Perceiving and Recognizing Automaton, 1957.
- [13] Neha Sharma, Vibhor Jain, and Anju Mishra. An Analysis of Convolutional Neural Networks for Image Classification. *Procedia Computer Science*, 132:377–384, 2018.
- [14] Connor Shorten and Taghi M. Khoshgoftaar. A survey on Image Data Augmentation for Deep Learning. *Journal of Big Data*, 6(1), 2019.
- [15] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *3rd International Conference on Learning Representations, ICLR 2015 - Conference Track Proceedings*, pages 1–14, 2015.
- [16] Sandro Skansi. *Introduction to Deep Learning: From Logical Calculus to Artificial Intelligence*. Springer, Cham, Switzerland, 2018.
- [17] Vladimir N. Vapnik. *The Nature of Statistical Learning Theory*. Springer, New York, NY, USA, 1995.

- [18] Dengsheng Zhang. *Fundamentals of Image Data Mining*. 2019.